

# Problem statement:

To predict how best the data fits and which model suits

## Linear Regression

### 1. Data Collection

In [3]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
```

In [4]:

```
df=pd.read_csv(r"C:\Users\chila\Downloads\insurance.csv")
df
```

Out[4]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

### 2. Data cleaning and preprocessing

In [5]:

```
df.head()
```

Out[5]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [6]:

```
df.tail()
```

Out[6]:

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

In [7]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [8]:

```
df.shape
```

Out[8]:

```
(1338, 7)
```

In [9]:

```
df.describe()
```

Out[9]:

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

## Finding null values

In [10]:

```
df.isnull().any()
```

Out[10]:

```
age      False
sex      False
bmi      False
children False
smoker   False
region   False
charges  False
dtype: bool
```

In [11]:

```
df.isnull().sum()
```

Out[11]:

```
age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

## Finding Duplicate values

In [12]:

```
df.duplicated().sum()
```

Out[12]:

1

In [13]:

```
df=df.drop_duplicates()  
df
```

Out[13]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1337 rows × 7 columns

In [28]:

```
T={"smoker":{"yes":1,"no":0}}  
df=df.replace(T)  
print(df)
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.92400
1	18	0	33.770	1	0	southeast	1725.55230
2	28	0	33.000	3	0	southeast	4449.46200
3	33	0	22.705	0	0	northwest	21984.47061
4	32	0	28.880	0	0	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	0	30.970	3	0	northwest	10600.54830
1334	18	1	31.920	0	0	northeast	2205.98080
1335	18	1	36.850	0	0	southeast	1629.83350
1336	21	1	25.800	0	0	southwest	2007.94500
1337	61	1	29.070	0	1	northwest	29141.36030

[1337 rows x 7 columns]

# Splitting the data into training data and testing data

In [30]:

```
x=df[['bmi']]
y=df['charges']
```


In [31]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=100)
```

## 3.Data visualization

In [32]:

```
plt.scatter(df['bmi'],df['charges'])
plt.show()
```



In [33]:

```
x.head(20)
```

Out[33]:

	bmi
0	27.900
1	33.770
2	33.000
3	22.705
4	28.880
5	25.740
6	33.440
7	27.740
8	29.830
9	25.840
10	26.220
11	26.290
12	34.400
13	39.820
14	42.130
15	24.600
16	30.780
17	23.845
18	40.300
19	35.300

In [34]:


```
y.head(15)
```

Out[34]:

```
0    16884.92400
1     1725.55230
2     4449.46200
3    21984.47061
4     3866.85520
5     3756.62160
6     8240.58960
7     7281.50560
8     6406.41070
9    28923.13692
10    2721.32080
11   27808.72510
12    1826.84300
13   11090.71780
14   39611.75770
Name: charges, dtype: float64
```

In [43]:

```
sns.lmplot(x='bmi',y='charges', order=2,data=df, ci=None)
plt.show()
```




In [46]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25, random_state=0)
lr=LinearRegression()
lr.fit(x_train,y_train)
print(lr.score (x_test,y_test))
```

0.060963613622574186

In [50]:


```
y_pred=lr.predict(x_test)
plt.scatter(x_test,y_test,color='c')
plt.plot(x_test,y_pred, color='r')
plt.show()
```





In [52]:

```
df500=df[:] [:500]
sns.lmplot(x='bmi',y='charges', order=2,ci=None, data=df500)
plt.show()
```



In [53]:

```
df500.fillna (method='ffill', inplace=True)
```

In [54]:

```
x=np.array(df500["bmi"]).reshape(-1,1)
y=np.array(df500['charges']).reshape(-1,1)
```

In [58]:

```
#Evaluation of model
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
lr=LinearRegression()
lr.fit(x_train,y_train)
y_pred=lr.predict(x_test)
r2=r2_score(y_test,y_pred)
print(r2)
```

0.060963613622574186

In [59]:

```
#The model accuracy is 6%.This is not the best suit model
```


## Ridge Regression

In [68]:

```
from sklearn.linear_model import Lasso, Ridge
I=df[['bmi','charges']]
plt.figure(figsize=(6,6))
sns.heatmap(I.corr(),annot=True)
```

Out[68]:

<Axes: >



In [69]:

```
features=df.columns [0:1]
target=df.columns[-1]
```

In [70]:

```
x=df[features].values
y=df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30, random_state=1)
print("The dimension of X_train is {}".format(x_train.shape))
print("The dimension of X_test is {}".format(x_test.shape))
```

The dimension of X\_train is (935, 1)  
The dimension of X\_test is (402, 1)

In [76]:

```
lr = LinearRegression()
lr.fit(x_train, y_train)
actual = y_test
train_score_lr = lr.score(x_train, y_train)
test_score_lr = lr.score(x_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.09099234134544743  
The test score for lr model is 0.07338609034045929

In [77]:

```
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(x_train, y_train)
test_score_ridge=ridgeReg.score(x_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.09099234107282062  
The test score for ridge model is 0.07338709056396597

## Lasso Regression

In [82]:

```
lasso= Lasso (alpha=10)
lasso.fit(x_train, y_train)
#train and test score for ridge regression
train_score_ls = lasso.score(x_train, y_train)
test_score_ls= lasso.score(x_test, y_test)
print("\nRidge Model:\n")
print("The train score for lasso model is {}".format(train_score_ls))
print("The test score for lasso model is {}".format(test_score_ls))
```

Ridge Model:

The train score for lasso model is 0.0909923379381713  
The test score for lasso model is 0.07338962361681955

## Logistic Regression

In [91]:

```
x=np.array(df['charges']).reshape(-1,1)
y=np.array(df['smoker']).reshape(-1,1)
df.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=1)
from sklearn.linear_model import LogisticRegression
lr=LogisticRegression(max_iter=10000)
```

In [92]:

```
lr.fit(x_train,y_train)
```

C:\Users\chila\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().  
y = column\_or\_1d(y, warn=True)

Out[92]:

▼	LogisticRegression
	LogisticRegression(max_iter=10000)


In [93]:

```
score=lr.score(x_test,y_test)
print(score)
```

0.9253731343283582

In [94]:

```
sns.regplot(x=x,y=y,data=df,logistic=True,ci=None)
plt.show()
```



In [108]:

```
#conclusion:We got the best fit curve for Logistic Regression
```

## Decision Tree

In [109]:

```
from sklearn.tree import DecisionTreeClassifier
clf=DecisionTreeClassifier(random_state=0)
clf.fit(x_train,y_train)
```

Out[109]:

```
DecisionTreeClassifier
DecisionTreeClassifier(random_state=0)
```

In [110]:

```
score=clf.score(x_test,y_test)
print(score)
```

```
0.900497512437811
```

# RANDOM FOREST

In [111]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

C:\Users\chila\AppData\Local\Temp\ipykernel\_4968\2210184639.py:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
rfc.fit(x_train,y_train)
```

Out[111]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In [112]:

```
params={'max_depth':[2,3,5,10,20],
        'min_samples_leaf':[5,10,20,50,100,200],
        'n_estimators':[10,25,30,50,100,200]}
```

In [113]:

```
from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rfc,param_grid=params,cv=2,scoring="accuracy")
```

In [114]:

```
grid_search.fit(x_train,y_train)
```

```
estimator.fit(x_train, y_train, **fit_params)
C:\Users\chila\AppData\Local\Programs\Python\Python310\lib\site-package
s\sklearn\model_selection\_validation.py:686: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change
the shape of y to (n_samples,), for example using ravel().
estimator.fit(X_train, y_train, **fit_params)
C:\Users\chila\AppData\Local\Programs\Python\Python310\lib\site-package
s\sklearn\model_selection\_validation.py:686: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change
the shape of y to (n_samples,), for example using ravel().
estimator.fit(X_train, y_train, **fit_params)
C:\Users\chila\AppData\Local\Programs\Python\Python310\lib\site-package
s\sklearn\model_selection\_validation.py:686: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change
the shape of y to (n_samples,), for example using ravel().
estimator.fit(X_train, y_train, **fit_params)
C:\Users\chila\AppData\Local\Programs\Python\Python310\lib\site-package
s\sklearn\model_selection\_validation.py:686: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change
the shape of y to (n_samples,), for example using ravel().
```

In [115]:

```
grid_search.best_score_
```

Out[115]:

0.9219284759969985

In [116]:

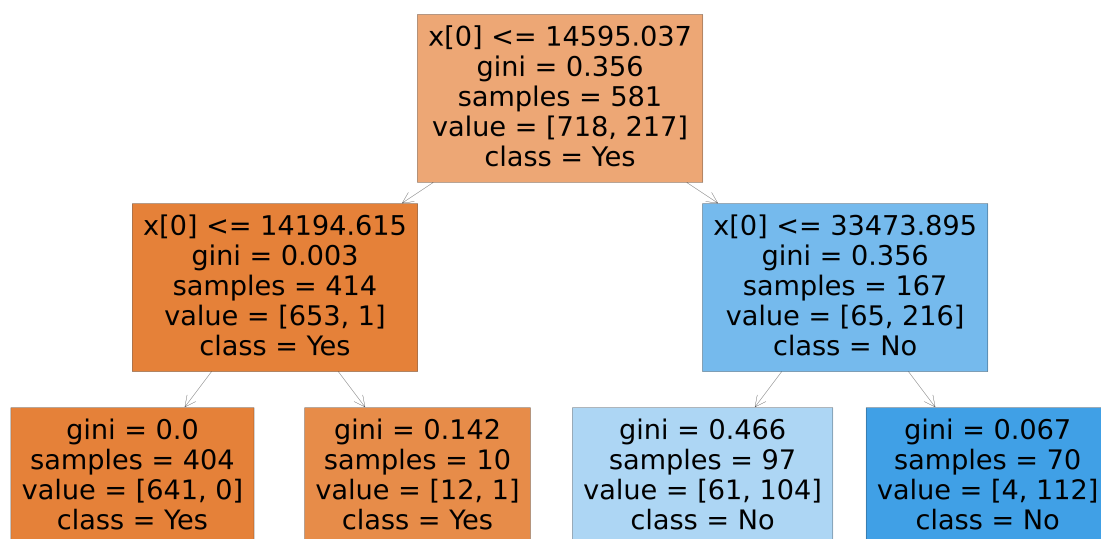
```
rf_best=grid_search.best_estimator_  
rf_best
```

Out[116]:

```
RandomForestClassifier  
RandomForestClassifier(max_depth=2, min_samples_leaf=10, n_estimators=10)
```

In [117]:

```
from sklearn.tree import plot_tree  
plt.figure(figsize=(80,40))  
plot_tree(rf_best.estimators_[5],class_names=['Yes','No'],filled=True);
```



In [118]:

```
score=rfc.score(x_test,y_test)  
print(score)
```

0.900497512437811

## Conclusion:

Finally we conclude that based on the accuracy of all models which we are implemented above the "Logistic Regression" is the best model.

